

# The Surprising Effectiveness of PPO in Cooperative Multi-Agent Games

## Paper

- <https://arxiv.org/pdf/2103.01955.pdf>

## Introduction

- Other multiagent algorithms: MADDPG, value decomposed q learning
- Revisits Proximal Policy Optimization (PPO) for multi agent settings

## Related work

- MARL algorithms
  - Centralized learning - directly learn single policy to produce joint actions of all agents
  - Decentralized learning - each agent optimizes its reward independently
  - Centralized Training and Decentralized execution - in between
    - Value decomposition methods

## Preliminaries

- Decentralized partially observable Markov decision processes (DEC-POMDP) with shared rewards
- $\langle S, A, O, R, P, n, \gamma \rangle$ 
  - $S$ : state space
  - $A$ : shared action space for each agent  $i$
  - $O_i = O(s; i)$ : local observation for each agent  $i$  at global state  $s$
  - $P(s'|s, A)$
  - $A = (a_1 \dots a_n)$
  - $R(s, A)$  shared reward function
  - $\gamma$ : discount factor  $r$
  - Agents use policy

## Mappo Details

---

### Algorithm 1 Recurrent-MAPPO

---

Initialize  $\theta$ , the parameters for policy  $\pi$  and  $\phi$ , the parameters for critic  $V$ , using Orthogonal initialization (Hu et al., 2020)  
Set learning rate  $\alpha$   
**while**  $step \leq step_{max}$  **do**  
  set data buffer  $D = \{\}$   
  **for**  $i = 1$  **to**  $batch\_size$  **do**  
     $\tau = []$  empty list  
    initialize  $h_{0,\pi}^{(1)}, \dots, h_{0,\pi}^{(n)}$  actor RNN states  
    initialize  $h_{0,V}^{(1)}, \dots, h_{0,V}^{(n)}$  critic RNN states  
    **for**  $t = 1$  **to**  $T$  **do**  
      **for all agents**  $a$  **do**  
         $p_t^{(a)}, h_{t,\pi}^{(a)} = \pi(o_t^{(a)}, h_{t-1,\pi}^{(a)}; \theta)$   
         $u_t^{(a)} \sim p_t^{(a)}$   
         $v_t^{(a)}, h_{t,V}^{(a)} = V(s_t^{(a)}, h_{t-1,V}^{(a)}; \phi)$   
      **end for**  
      Execute actions  $u_t$ , observe  $r_t, s_{t+1}, o_{t+1}$   
       $\tau \leftarrow [s_t, o_t, h_{t,\pi}, h_{t,V}, u_t, r_t, s_{t+1}, o_{t+1}]$   
    **end for**  
    Compute advantage estimate  $\hat{A}$  via GAE on  $\tau$ , using PopArt  
    Compute reward-to-go  $\hat{R}$  on  $\tau$  and normalize with PopArt  
    Split trajectory  $\tau$  into chunks of length  $L$   
    **for**  $l = 0, 1, \dots, T//L$  **do**  
       $D = D \cup (\tau[l : l + T], \hat{A}[l : l + L], \hat{R}[l : l + L])$   
    **end for**  
  **end for**  
  **for** mini-batch  $k = 1, \dots, K$  **do**  
     $b \leftarrow$  random mini-batch from  $D$  with all agent data  
    **for each data chunk**  $c$  in the mini-batch  $b$  **do**  
      update RNN hidden states for  $\pi$  and  $V$  from first hidden state in data chunk  
    **end for**  
  **end for**  
  Adam update  $\theta$  on  $L(\theta)$  with data  $b$   
  Adam update  $\phi$  on  $L(\phi)$  with data  $b$   
**end while**

---

- Mappo trains 2 separate neural networks
  - Actor network with parameters  $\theta$
  - Value function network (critic) with parameters  $\phi$

The actor network is trained to maximize the objective

$$L(\theta) = [\frac{1}{Bn} \sum_{i=1}^B \sum_{k=1}^n \min(r_{\theta,i}^{(k)} A_i^{(k)}, \text{clip}(r_{\theta,i}^{(k)}, 1 - \epsilon, 1 + \epsilon) A_i^{(k)})] + \sigma \frac{1}{Bn} \sum_{i=1}^B \sum_{k=1}^n S[\pi_{\theta}(o_i^{(k)})], \text{ where}$$

$r_{\theta,i}^{(k)} = \frac{\pi_{\theta}(a_i^{(k)} | o_i^{(k)})}{\pi_{\theta_{old}}(a_i^{(k)} | o_i^{(k)})}$ .  $A_i^{(k)}$  is computed using the GAE method,  $S$  is the policy entropy, and  $\sigma$  is the entropy coefficient hyperparameter.

•

The critic network is trained to minimize the loss function

$$L(\phi) = \frac{1}{Bn} \sum_{i=1}^B \sum_{k=1}^n (\max[(V_{\phi}(s_i^{(k)}) - \hat{R}_i)^2, (\text{clip}(V_{\phi}(s_i^{(k)}), V_{\phi_{old}}(s_i^{(k)}) - \varepsilon, V_{\phi_{old}}(s_i^{(k)}) + \varepsilon) - \hat{R}_i)^2],$$

where  $\hat{R}_i$  is the discounted reward-to-go.

In the loss functions above,  $B$  refers to the batch size and  $n$  refers to the number of agents.

If the critic and actor networks are RNNs, then the loss functions additionally sum over time, and the networks are trained via Backpropagation Through Time (BPTT). Pseudocode for recurrent-MAPPO is shown in Alg. 1.

- 
- Trained via backpropagation through time  
[https://d2l.ai/chapter\\_recurrent-neural-networks/bptt.html](https://d2l.ai/chapter_recurrent-neural-networks/bptt.html)
- Generalized Advantage Estimation
  -

## Factors influential to ppo performance

- Value normalization
- Input representation to value function
  - “the fundamental difference between many multi-agent CTDE PG algorithms and fully decentralized PG methods is the input to the value network”
  - Options
    - Concatenation Of local observations
    - Environment provided global state
    - Agent specific global state
    - Feature pruned agent specific global state
- PPO clipping
  - PPO clipped importance ratio and value loss
  - Larger clipping values means larger updates to policy and value function ‘
  - Small clipping values → agent’s policy less likely to change per episode
- Ppo batch size
  - Shouldn’t train it for too long

## Parameter sharing

- All agents share the same networks
- Recurrent data chunk length

## Death masking

- When an agent dies these agent specific features become zero

<https://arxiv.org/abs/2304.09870>

Main takeaways